

Klute Analysis Report for Twitter Project

Software/hardware used:

Server: MySQL (and MySQL Workbench)
Storage Engine:
Language: Python (in Pycharm)
Database: Terminal
OS: Mac Pro
Processor: M1
RAM: 16GB
Storage Engine: InnoDB
Prepared Statements: Yes
Secondary Indexing: No

Set Up:

Used environmental variables (inputted placeholders for TA)
Set up serve connection with MySQL Workbench, specified in advanced
OPT_LOCAL_INFILE=1 to load follows.csv into FOLLOWS table

Approach:

I set up my FOLLOWS table by reading in the csv file in my main function and using the api function I created called load_data_from_csv (which is in the tweets_mysql.py file).

I then created a dataframe of df_TWEET and looped through each row in the df and using the Tweet object and use the api function of post_tweet to insert_one through dbutils into the TWEET table in the database. This loads the 1,000,000 tweets into the table. I used a timer function to calculate the average tweets per second which came to 3779.46.

I then created api functions that get a random user, calculate the timelines per second through calling on the home_timeline function for the random user which I pulled by creating a get_random_user_ids function (in tweets_mysql.py file).

I tested it with differing amounts of users, and landed on 500 as it took a really long time to run more than that. My outputted number of timelines retrieved per second: 2.47.

I also created get_followers, get_followees, get tweets to use for future implementation.

Final outputs:

```
/Users/sarahklute/anaconda3/envs/ds/bin/python
/Users/sarahklute/Documents/northeastern/year2/ds4300/hw1/tweets_tester.py
root
Data inserted successfully.
Number of tweets processed: 1000000
```

Time taken to process tweets: 264.59 seconds

Average Tweets per Second: 3779.46

Sample Timeline for User 4857:

	tweet_id	user_id	tweet_text \
0	2009927	8542	y g erem knvcgjrqnnyz v ojeymb kpqhwex pvuw o ...
1	2008972	8542	hfty bljb edfaj qp n ld uob v euzne xoy...
2	2004054	2604	h wm axdhjnyixaslel wa upfabe hqow qxe o rtvj...
3	2004951	6142	yvp iqyfb jas
4	1999029	8542	c fbmbay ase oi lcuiyybdfu h n hgeeyebujbg m e...
5	1995189	6142	yy kwr vfhgol pfryundnprugjrs etcn v k
6	1992183	8542	qvwbnwak e uelsagqly qqsjiwux saanymfq dyaxz ...
7	1993150	8542	k wnbhuybsnqvqvqreija xramnrhqgtfy
8	1992413	7520	n lr txwpoie cvwg edwr bhurmvska mat jhmw
9	1990870	8542	n spskyr tbvq wy nee cx j ar cxcyirp mrsli...

	tweet_ts
0	2024-01-21 18:43:53
1	2024-01-21 18:43:53
2	2024-01-21 18:43:52
3	2024-01-21 18:43:52
4	2024-01-21 18:43:51
5	2024-01-21 18:43:50
6	2024-01-21 18:43:49
7	2024-01-21 18:43:49
8	2024-01-21 18:43:49
9	2024-01-21 18:43:48

Number of timelines retrieved per second: 2.47

Time taken to retrieve timelines for 500 users: 202.66 seconds

Process finished with exit code 0

Possible Limitations:

- 16GB RAM on M1 might not be the most efficient processor.
- Mac OS, some libraries are optimized for Windows, Linux
- Random user calls within other function might slow down the timeline for fetching the home timeline.
- Absence of secondary indexing might be impacting query performance.
- Dependency on anaconda environment
 - o I set up environmental variables.